

M07 Homework

Michael Vaden, mtv2eva

```
In [ ]: import pandas as pd
import numpy as np
from glob import glob
import re
import nltk
import plotly_express as px
from lib.textparser import TextParser
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer, TfidfTransformer

from numpy.linalg import norm
from scipy.spatial.distance import pdist
import scipy.cluster.hierarchy as sch
import matplotlib.pyplot as plt
from scipy.linalg import eigh
from sklearn.decomposition import PCA
import seaborn as sns
```

```
In [ ]: import configparser
config = configparser.ConfigParser()
config.read("../env.ini")
data_home = config['DEFAULT']['data_home']
output_dir = config['DEFAULT']['output_dir']
data_prefix = 'novels'
```

Import the LIB and CORPUS tables from the provided CSV files. The index OHCO for CORPUS is ['book_id','chapter_id','para_num','sent_num','token_num'] The index for LIB is 'book_id'.

```
In [ ]: OHCO = ['book_id', 'chap_id', 'para_num', 'sent_num', 'token_num']
```

```
In [ ]: LIB = pd.read_csv(f"{data_home}/{data_prefix}-LIB.csv").set_index('book_id')
CORPUS = pd.read_csv(f"{data_home}/{data_prefix}-CORPUS.csv").set_index(OHCO)
```

```
In [ ]: LIB
```

Out []:

	genre_id	author_id
book_id		
secretadversary	d	christie
styles	d	christie
moonstone	d	collins
adventures	d	doyle
baskervilles	d	doyle
scarlet	d	doyle
signoffour	d	doyle
marieroget	d	poe
ruemorgue	d	poe
northangerabbey	g	austen
christmascarole	g	dickens
monk	g	lewis
pitandpendulum	g	poe
reddeath	g	poe
usher	g	poe
udolpho	g	radcliffe
oldenglishbaron	g	reeve
frankenstein	g	shelley
dracula	g	stoker
castleofotranto	g	walpole

```
In [ ]: CORPUS
```

```
Out [ ]:
```

	book_id	chap_id	para_num	sent_num	token_num	pos	term_str
	secretadversary	1	0	1	0	DT	the
					1	NNP	young
					2	NNP	adventurers
					3	NNP	ltd
			1	0	0	JJ	tommy

	baskervilles	11	114	1	7	RBR	more
					8	JJ	comfortable
					9	IN	outside
					10	IN	than
					11	NN	in

1500417 rows × 2 columns

Extract a VOCAB table and add max_pos as a feature.

```
In [ ]:
```

```
CORPUS = CORPUS[CORPUS.term_str != '']
CORPUS['pos_group'] = CORPUS.pos.str[:2]

VOCAB = CORPUS.term_str.value_counts().to_frame('n').sort_index()
VOCAB.index.name = 'term_str'
VOCAB['n_chars'] = VOCAB.index.str.len()
VOCAB['p'] = VOCAB.n / VOCAB.n.sum()
VOCAB['i'] = -np.log2(VOCAB.p)

VOCAB['max_pos'] = CORPUS[['term_str', 'pos']].value_counts().unstack(fill_value=0).idxmax(1)
```

```
In [ ]:
```

VOCAB

```
Out [ ]:
```

	n	n_chars	p	i	max_pos
term_str					
a	28533	1	1.901674e-02	5.716586	DT
aback	9	5	5.998340e-06	17.347005	NN
abaft	2	5	1.332965e-06	19.516930	IN
abandon	44	7	2.932522e-05	15.057499	VB
abandoned	68	9	4.532079e-05	14.429467	VTB
...
à	3	1	1.999447e-06	18.931968	NNP
æt	1	2	6.664823e-07	20.516930	NN
ætæt	1	4	6.664823e-07	20.516930	VBD
ça	2	2	1.332965e-06	19.516930	JJ
émeutes	1	7	6.664823e-07	20.516930	PDT

27396 rows × 5 columns

Compute TFIDF and VOCAB['dfidf'] for the CORPUS using the following parameters:

```
bag = ['book_id', 'chap_id']
```

```
tf_method = 'max'
```

```
idf_method = 'standard'
```

```
In [ ]:
```

```
def create_bag_of_words(CORPUS, bag):
    BOW = CORPUS.groupby(bag+['term_str']).term_str.count().to_frame('n')
    return BOW

idf_method = 'standard'

def get_TFIDF(BOW, tf_method):
    DTCM = BOW.n.unstack(fill_value=0)

    DF = DTCM.astype('bool').sum()
    N = len(DTCM)

    if tf_method == 'sum':
        TF = DTCM.T / DTCM.T.sum()

    elif tf_method == 'max':
        TF = DTCM.T / DTCM.T.max()

    elif tf_method == 'log':
        TF = np.log2(1 + DTCM.T)

    elif tf_method == 'raw':
        TF = DTCM.T
```

```

elif tf_method == 'double_norm':
    TF = DTCM.T / DTCM.T.max()

elif tf_method == 'binary':
    TF = DTCM.T.astype('bool').astype('int')

TF = TF.T

if idf_method == 'standard':
    IDF = np.log2(N / DF)

elif idf_method == 'max':
    IDF = np.log2(DF.max() / DF)

elif idf_method == 'smooth':
    IDF = np.log2((1 + N) / (1 + DF)) + 1

return TF * IDF, DF * IDF

```

```

In [ ]: TFIDF, DFIDF = get_TFIDF(create_bag_of_words(CORPUS, ['book_id', 'chap_id']), 'max')

TFIDF

```

```

Out [ ]:
      term_str  a  aback  abaft  abandon  abandoned  abandoning  abandons  abasement  abashed  abate  ...  zoöphagy  zufalle  zum  zuniga  zusammen  à  æ
book_id  chap_id
adventures  1  0.0    0.0    0.0  0.00000  0.006493    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           2  0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           3  0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           4  0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           5  0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           ...  ...    ...    ...    ...    ...    ...    ...    ...    ...    ...  ...    ...    ...    ...    ...    ...  ...  .
udolpho     54 0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           55 0.0    0.0    0.0  0.00000  0.009341    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           56 0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
           57 0.0    0.0    0.0  0.00000  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.
usher        1 0.0    0.0    0.0  0.01253  0.000000    0.0    0.0    0.0    0.0    0.0  ...    0.0    0.0    0.0    0.0    0.0  0.0  0.

```

320 rows × 27396 columns

Create a DOC table from the TFIDF index in which each row represents a bag, i.e. a chapter. In other words, it should have ['book_id', 'chap_id'] as its index. This table should have information from the LIB table added to it, so that each chapter is identified with an author, title, and genre. These data will appear in your visualizations. For example, in a scatter plot of documents in the first two principle components, you will want to know the book and chapter of each data point.

```

In [ ]: DOC = TFIDF.reset_index()[['book_id', 'chap_id']]
DOC = DOC.merge(LIB, on='book_id', how='left')
DOC.set_index(['book_id', 'chap_id'], inplace=True)
DOC

```

```

Out [ ]:
      genre_id  author_id
book_id  chap_id
adventures  1      d      doyle
           2      d      doyle
           3      d      doyle
           4      d      doyle
           5      d      doyle
           ...  ...      ...      ...
udolpho     54      g      radcliffe
           55      g      radcliffe
           56      g      radcliffe
           57      g      radcliffe
usher        1      g      poe

```

320 rows × 2 columns

Create a reduced version of the TFIDF table with only the top 1000 nouns (i.e. NN and NNS) in descending order of DFIDF.

Do not "collapse" table -- keep the index as (book_id, chap_id).

```

In [ ]: nouns = ['NN', 'NNS']
top1000 = VOCAB.merge(DFIDF.to_frame('DFIDF'), on='term_str').query("max_pos in @nouns").sort_values('DFIDF', ascending=False).head(1000).index
TFIDF = TFIDF[[col for col in top1000]]
TFIDF

```

```
Out [ ]:
```

	term_str	yours	reply	order	curiosity	memory	company	feelings	opportunity	book	spirit	...	humanity	rank	contempt	apprehension
	book_id	chap_id														
adventures	1	0.006454	0.000000	0.003227	0.003227	0.006400	0.006400	0.000000	0.000000	0.003282	0.003282	...	0.000000	0.000000	0.000000	0
	2	0.009346	0.000000	0.009346	0.000000	0.006178	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.000000	0
	3	0.004089	0.008178	0.000000	0.000000	0.004054	0.008109	0.000000	0.000000	0.004159	0.000000	...	0.000000	0.000000	0.000000	0
	4	0.002721	0.002721	0.005442	0.000000	0.002698	0.000000	0.000000	0.002767	0.000000	0.000000	...	0.000000	0.000000	0.000000	0
	5	0.003043	0.000000	0.003043	0.003043	0.003017	0.000000	0.000000	0.000000	0.006190	0.000000	...	0.000000	0.000000	0.000000	0
...
udolpho	54	0.007676	0.003838	0.003838	0.003838	0.003806	0.003806	0.007611	0.007808	0.000000	0.000000	...	0.017027	0.000000	0.000000	0
	55	0.000000	0.000000	0.000000	0.004643	0.000000	0.004604	0.000000	0.000000	0.000000	0.004722	...	0.000000	0.000000	0.010299	0
	56	0.000000	0.007822	0.000000	0.000000	0.000000	0.015512	0.000000	0.000000	0.000000	0.015912	...	0.000000	0.000000	0.017351	0
	57	0.000000	0.000000	0.000000	0.000000	0.033189	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.000000	0
usher	1	0.000000	0.002566	0.002566	0.000000	0.002544	0.000000	0.005088	0.002610	0.005219	0.015657	...	0.005691	0.011382	0.000000	0

320 rows x 1000 columns

Write a function that computes PCA from a given document-term count matrix (this included weighted counts, such a tfidf). It should return three dataframes: LOADINGS (the term-component matrix), DCM (the document-component matrix), COMPINF (the component information table). Give it the following parameters:

X # The input matrix

k # The number of components to generate

norm_docs # True or False

center_by_mean # True or False

center_by_variance # True or False

Compute PCA from the feature-reduced TFIDF table using your function. Use the following parameter values: X = TFIDF_REDUCED Or whatever you called your reduced TFIDF table k=10 norm_docs = True center_by_mean = False center_by_variance = False Visualize your results using scatter plots and box plots.

```
In [ ]:
```

```
def get_PCA(X, k, norm_docs = True, center_by_mean = True, center_by_variance = True):
    if norm_docs:
        X = (X.T / np.linalg.norm(X, 2, axis=1)).T

    if center_by_mean:
        X = X - X.mean(axis=0)

    if center_by_variance:
        X = X / X.std(axis=0)

    pca = PCA(n_components=k)
    pca.fit(X)

    loadings = pd.DataFrame(pca.components_.T, index=X.columns)
    loadings.columns = [f'PC{i+1}' for i in range(k)]

    dcm = pd.DataFrame(pca.transform(X), index=X.index)
    dcm.columns = [f'PC{i+1}' for i in range(k)]

    comp = pd.DataFrame({
        'EVR': pca.explained_variance_ratio_,
        'CEVR': np.cumsum(pca.explained_variance_ratio_)
    }, index=[f'PC{i+1}' for i in range(k)])

    return loadings, dcm, comp

LOADINGS, DCM, COMPINF = get_PCA(X=TFIDF, k=10, norm_docs=True, center_by_mean=False, center_by_variance=False)
DCM = DOC.join(DCM)
```

```
In [ ]:
```

DCM

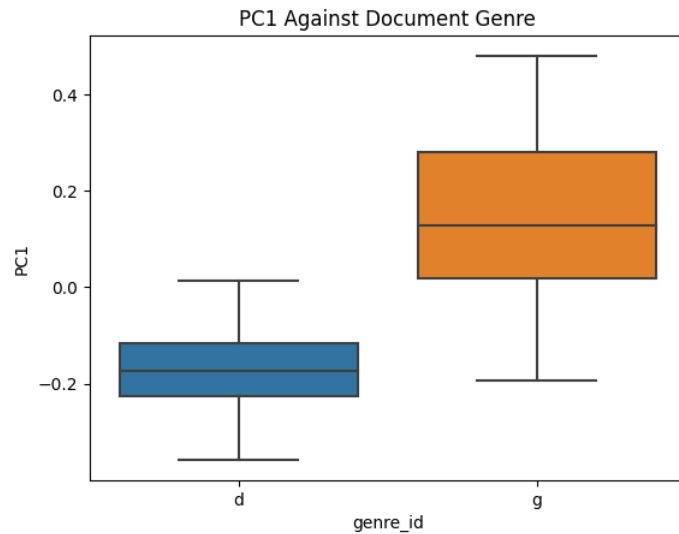
```
Out [ ]:
```

	genre_id	author_id	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10	
	book_id	chap_id											
adventures	1	d	doyle	-0.237319	0.038768	0.092062	0.029442	-0.191567	0.093441	-0.118399	0.011765	-0.021905	-0.028254
	2	d	doyle	-0.295389	0.074901	0.035613	0.001857	-0.204501	0.066934	-0.111104	-0.109326	-0.020459	0.026691
	3	d	doyle	-0.291923	-0.074162	0.035053	0.008414	-0.149578	0.112748	-0.103293	-0.041946	0.000359	0.040192
	4	d	doyle	-0.162984	-0.025831	0.026673	0.081026	-0.141952	0.057792	-0.204866	-0.071068	0.158094	0.085781
	5	d	doyle	-0.186560	0.049075	0.012450	0.066207	-0.145323	0.097464	-0.043226	-0.144036	0.006083	0.152767
...
udolpho	54	g	radcliffe	0.379714	-0.132359	-0.030655	-0.065428	0.040511	0.206042	-0.063613	-0.036247	0.028850	-0.000011
	55	g	radcliffe	0.303586	-0.185496	0.050914	0.036115	0.065106	0.181303	-0.023661	-0.113286	0.074211	-0.011379
	56	g	radcliffe	0.391705	0.145064	-0.276399	-0.235498	-0.109495	0.141362	0.046051	0.030292	-0.016949	0.121353
	57	g	radcliffe	0.286226	0.017426	-0.042757	-0.112652	0.019443	0.009336	-0.095355	-0.012079	-0.051540	0.123404
usher	1	g	poe	0.132124	0.119597	0.032166	0.126673	-0.018076	0.044846	-0.036131	-0.058944	-0.032415	-0.193522

320 rows x 12 columns

1. Looking at the documents plotted against the first principle component (PC), which genre has the more narrow range, i.e. distance between the minimum and maximum values? This can be seen using a box plot

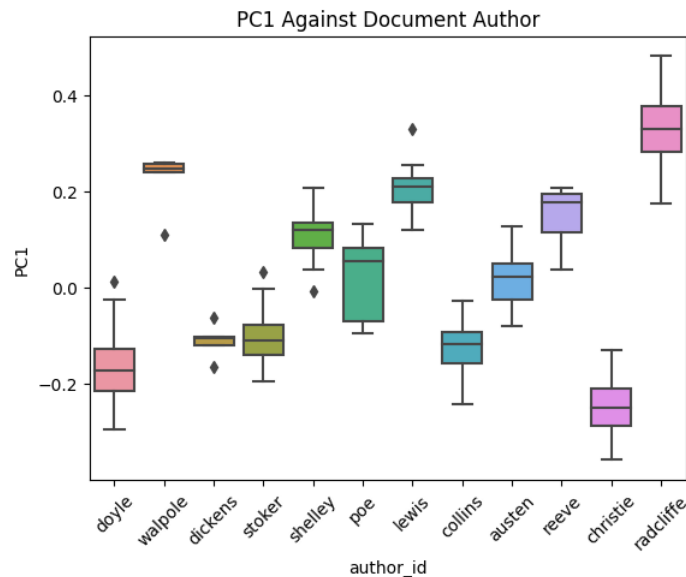
```
In [ ]: sns.boxplot(data=DCM, x = 'genre_id', y = 'PC1')
plt.title('PC1 Against Document Genre')
plt.show()
```



We can see from the boxplot above for the first principle component that the **d (detective)** genre has a more narrow range between the minimum and maximum values for the first principle component.

2. Looking at the documents plotted against the first PC, which author has the highest absolute value, in terms of both mean and range? In other words, which author is farthest from 0? Again, the box plots of each author are useful here.

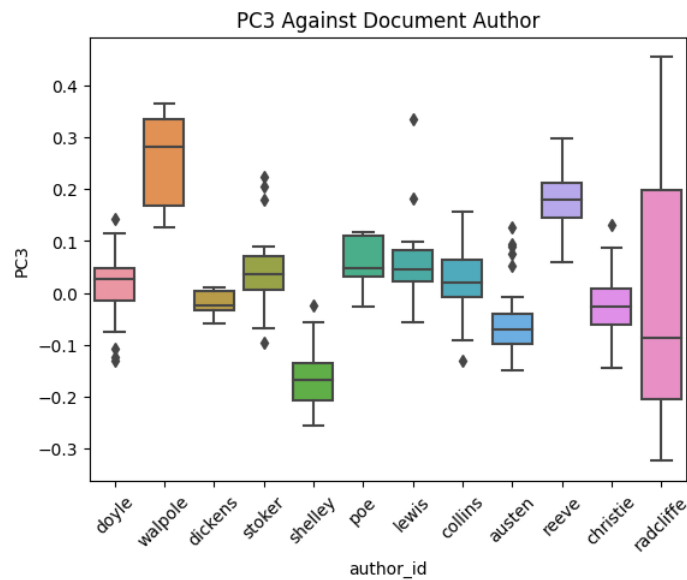
```
In [ ]: sns.boxplot(data=DCM, x = 'author_id', y = 'PC1')
plt.title('PC1 Against Document Author')
plt.xticks(rotation=45)
plt.show()
```



We can see from the boxplot above for the first principle component that the author **Radcliffe** has both the highest absolute mean value of mean and range.

3. In the third PC, which author has, by far, the maximum range?

```
In [ ]: sns.boxplot(data=DCM, x = 'author_id', y = 'PC3')
plt.title('PC3 Against Document Author')
plt.xticks(rotation=45)
plt.show()
```



We can see from the boxplot above for the third principle component that the author **Radcliffe** has the highest maximum range by far

4. Looking at the loadings for the second PC, how would you characterize the opposition, based on the top three words at each pole?

```
In [ ]: LOADINGS['PC2'].sort_values(ascending=False).to_frame().head(3)
```

```
Out [ ]:      PC2
term_str
mountains  0.176642
woods      0.140543
sea        0.103898
```

```
In [ ]: LOADINGS['PC2'].sort_values(ascending=True).to_frame().head(3)
```

```
Out [ ]:      PC2
term_str
brother   -0.153843
engagement -0.115630
father    -0.113204
```

For the loadings of the second principle component, we can see that there is the strongest positive terms are related to nature. In contrast, we can see that the strongest negative terms are all related to family. We can see a potential contrast here for the opposition between nature and family

5. Recompute the principle components with center_by_variance set to True. This will change the words that appear at the extremes of the first PC. Does this change your interpretation in the previous question?

```
In [ ]: LOADINGS5, DCM5, COMPINF5 = get_PCA(X=TFIDF, k=10, norm_docs=True, center_by_mean=False, center_by_variance=True)
DCM5 = DOC.join(DCM5)
```

```
In [ ]: LOADINGS5['PC1'].sort_values(ascending=False).to_frame().head(3)
```

```
Out [ ]:      PC1
term_str
countenance 0.095031
length      0.094551
melancholy  0.087428
```

```
In [ ]: LOADINGS5['PC1'].sort_values(ascending=True).to_frame().head(3)
```

```
Out [ ]:      PC1
term_str
matter     -0.085165
anything    -0.082012
course     -0.075488
```

instructions unclear, recomputing PCA and examining extremes at first PC. Compared to the previous question, we do not see clear themes of family or nature. Rather, the extremes of the first PC do not have as much of an association with each other that we can discern when we are centered by variance.